

S1-24_AIMLCZG523 – MLOps - Assignment 1

Group No. 120

<i>Name</i>	<i>StudentID</i>	<i>Contribution</i>
JOSHI NIRANJAN SURYAKANT	2023AC05011	100%
PRATEEK RALHAN	2023AC05673	100%
KESHARKAR SURAJ SANJAY	2023AD05004	100%
SAURABH SUNIT JOTSHI	2023AC05565	100%

MLOps Architecture Summary - California Housing Price Prediction Pipeline

Project Overview

End-to-end MLOps pipeline for California Housing Price Prediction using **Linear Regression and Decision Tree models, implementing industry best practices for model development, deployment, and monitoring.**

Architecture Components

- Data Management & Versioning**
 - Dataset*: California Housing dataset (housing.csv) with 8 features (MedInc, HouseAge, AveRooms, AveBedrms, Population, AveOccup, Latitude, Longitude)
 - Data Versioning*: DVC (Data Version Control) for tracking dataset changes
 - Data Pipeline*: Automated data loading via `src/utlis.py` using scikit-learn's `fetch_california_housing`
 - Storage*: Local data directory with DVC tracking files
- Model Development & Training**
 - Algorithms*: Multiple model comparison (Linear Regression, Decision Tree, Ridge Regression)
 - Configuration*: YAML-based config (`src/config.yaml`) with hyperparameters (alpha, test_size, random_state, max_depth)
 - Training Pipeline*: `src/train.py` with automated model selection based on MSE
 - Model Storage*: Joblib serialization to `models/ridge_model.pkl`
 - Code Quality*: Pre-commit hooks (Black, Flake8, Prettier) for code formatting and linting
- Experiment Tracking & Model Registry**
 - Platform*: MLflow for experiment tracking and model versioning
 - Tracking*: Automatic logging of parameters, metrics (MSE), and model artifacts
 - Model Registry*: Registered models with versioning (`BestHousingModel`)
 - Storage*: SQLite backend (`mlflow_data/mlflow.db`) and local artifact storage (`mlruns/`)
 - UI*: MLflow tracking server on port 5555
- API development & Deployment**
 - Framework*: FastAPI with automatic OpenAPI documentation
 - Endpoints*:
 - `/predict` - Housing price predictions
 - `/metrics` - Prometheus metrics exposure




- iii. ``/retrain`` - Model retraining trigger
 - c. *Input Validation*: Pydantic models for request validation
 - d. *Containerization*: Docker with multi-stage build process
 - e. *Docker Hub*: Published image (``niranjanjoshi14/housing-api:latest``)
5. Monitoring & Observability
- a. *Metrics Collection*: Prometheus client integration with custom metrics
 - ``prediction_requests_total`` - Request counter
 - ``prediction_latency_seconds`` - Response time histogram
 - b. *Database Logging*: SQLite database (``prediction_logs.db``) for prediction history
 - c. *Monitoring Stack*:
 - Prometheus (port 9090) for metrics scraping
 - Grafana (port 3000) for visualization dashboards
 - d. *Logging*: Structured logging with timestamps and input/output data
6. CI/CD Pipelines
- a. *Platform*: GitHub Actions with automated workflows
 - b. *Triggers*: Push to main branch, pull requests, data changes
 - c. *Pipeline Stages*:
 - 1. Code checkout and Python setup
 - 2. Dependency installation and caching
 - 3. Code quality checks (Flake8, Black)
 - 4. Docker image building and publishing
 - 5. Model training validation
 - 6. Automated testing
 - d. *Security*: Docker Hub authentication via secrets
7. Infrastructure & Orchestration
- a. *Container Orchestration*: Docker Compose with 4 services
 - API service (FastAPI)
 - Prometheus (metrics collection)
 - Grafana (visualization)
 - MLflow (experiment tracking)
 - b. *Networking*: Internal service communication with external port mapping
 - c. *Volumes*: Persistent storage for Grafana and MLflow data
 - d. *Environment*: Python 3.10 with all dependencies in requirements.txt
8. Data Pipeline Automation
- a. *DVC Pipeline*: Automated training pipeline (``dvc.yaml``)
 - b. *Dependencies*: Clear dependency tracking between data, code, and models
 - c. *Reproducibility*: Deterministic training with fixed random seeds
 - d. *Artifact Management*: Automatic model artifact generation and tracking

Key Metrics & KPIs

- *Model Performance*: MSE (Mean Squared Error) tracking
- *API Performance*: Request latency, throughput, error rates
- *System Health*: Container status, resource utilization
- *Data Quality*: Input validation success rates
- *Deployment Success*: CI/CD pipeline success rates

This architecture demonstrates a **production-ready MLOps pipeline with comprehensive monitoring, automated deployment, and scalable infrastructure following industry best practices.**

Appendix:

- Github repository : [link](#) 
- Video demo: [link](#) 
- Published docker image: [link](#) 
- System Architecture (mermaid diagram): [link](#) 